

```
In [16]: import numpy as np  
import pandas as pd
```

```
In [17]: #Interactive plots using plotly  
import plotly.graph_objects as go  
import plotly.offline as plt  
import plotly.figure_factory as ff
```

```
In [18]: #For Model training  
from sklearn.model_selection import train_test_split  
from sklearn.ensemble import RandomForestClassifier  
from sklearn.metrics import confusion_matrix, classification_report  
from imblearn.over_sampling import SMOTE  
from imblearn.under_sampling import NearMiss
```

```
In [19]: #Prepare data  
source = pd.read_csv(r'C:\Users\zishe\spyder-py3\repos\CIS 568\data.csv')
```

```
In [20]: #Check for null values  
source.isnull().any()
```

```
Out[20]: Age           False  
Attrition      False  
BusinessTravel  False  
DailyRate       False  
Department     False  
DistanceFromHome False  
Education       False  
EducationField  False  
EmployeeCount   False  
EmployeeNumber  False  
EnvironmentSatisfaction False  
Gender          False  
HourlyRate      False  
JobInvolvement  False  
JobLevel        False  
JobRole         False  
JobSatisfaction False  
MaritalStatus   False  
MonthlyIncome    False  
MonthlyRate     False  
NumCompaniesWorked False  
Over18          False  
OverTime         False  
PercentSalaryHike False  
PerformanceRating False  
RelationshipSatisfaction False  
StandardHours   False  
StockOptionLevel False  
TotalWorkingYears False  
TrainingTimesLastYear False  
WorkLifeBalance  False  
YearsAtCompany  False  
YearsInCurrentRole False  
YearsSinceLastPromotion False  
YearsWithCurrManager False  
dtype: bool
```

```
In [21]: #What kind of dtypes are there
```

```
source.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 35 columns):
Age           1470 non-null int64
Attrition     1470 non-null object
BusinessTravel 1470 non-null object
DailyRate      1470 non-null int64
Department    1470 non-null object
DistanceFromHome 1470 non-null int64
Education      1470 non-null int64
EducationField 1470 non-null object
EmployeeCount   1470 non-null int64
EmployeeNumber  1470 non-null int64
EnvironmentSatisfaction 1470 non-null int64
Gender         1470 non-null object
HourlyRate     1470 non-null int64
JobInvolvement 1470 non-null int64
JobLevel       1470 non-null int64
JobRole        1470 non-null object
JobSatisfaction 1470 non-null int64
MaritalStatus   1470 non-null object
MonthlyIncome   1470 non-null int64
MonthlyRate     1470 non-null int64
NumCompaniesWorked 1470 non-null int64
Over18          1470 non-null object
OverTime         1470 non-null object
PercentSalaryHike 1470 non-null int64
PerformanceRating 1470 non-null int64
RelationshipSatisfaction 1470 non-null int64
StandardHours   1470 non-null int64
StockOptionLevel 1470 non-null int64
TotalWorkingYears 1470 non-null int64
TrainingTimesLastYear 1470 non-null int64
WorkLifeBalance 1470 non-null int64
YearsAtCompany   1470 non-null int64
YearsInCurrentRole 1470 non-null int64
YearsSinceLastPromotion 1470 non-null int64
YearsWithCurrManager 1470 non-null int64
dtypes: int64(26), object(9)
memory usage: 402.0+ KB
```

```
In [22]: #Check the number of unique values in the data
```

```
source.nunique()
```

```
Out[22]: Age            43
Attrition       2
BusinessTravel  3
DailyRate        886
Department      3
DistanceFromHome 29
Education        5
EducationField   6
EmployeeCount    1
EmployeeNumber   1470
EnvironmentSatisfaction 4
Gender           2
HourlyRate       71
JobInvolvement   4
JobLevel          5
JobRole           9
JobSatisfaction  4
MaritalStatus     3
MonthlyIncome     1349
MonthlyRate       1427
NumCompaniesWorked 10
Over18            1
OverTime          2
PercentSalaryHike 15
PerformanceRating 2
RelationshipSatisfaction 4
StandardHours     1
StockOptionLevel  4
TotalWorkingYears 40
TrainingTimesLastYear 7
WorkLifeBalance   4
YearsAtCompany    37
YearsInCurrentRole 19
YearsSinceLastPromotion 16
YearsWithCurrManager 18
dtype: int64
```

```
In [23]: #Dropping unnecessary data
```

```
features = source.drop(columns=['Attrition', 'EmployeeCount', 'Over18', 'StandardHours', 'EmployeeNumber'])
```

```
In [24]: #Label OHE
label = source['Attrition']
label = label.apply(lambda x:{'Yes': 1, 'No': 0}[x])
label = label.to_frame()
```

```
In [25]: #Split numerical and categorical features
numFeatures = features.select_dtypes(include=['int64'])
catFeatures = features.select_dtypes(include=['object'])
#OHE categorical features
catFeatures = pd.get_dummies(catFeatures)
```

```
In [26]: data = pd.concat([label, numFeatures, catFeatures], axis=1)
```

```
In [27]: data.head()
```

Out[27]:

	Attrition	Age	DailyRate	DistanceFromHome	Education	EnvironmentSatisfaction	HourlyRate	JobInvolvement	JobLevel	JobSatisfaction	...	JobRole_Man
0	1	41	1102		1	2	2	94	3	2	4	...
1	0	49	279		8	1	3	61	2	2	2	...
2	1	37	1373		2	2	4	92	2	1	3	...
3	0	33	1392		3	4	4	56	3	1	3	...
4	0	27	591		2	1	1	40	3	1	2	...

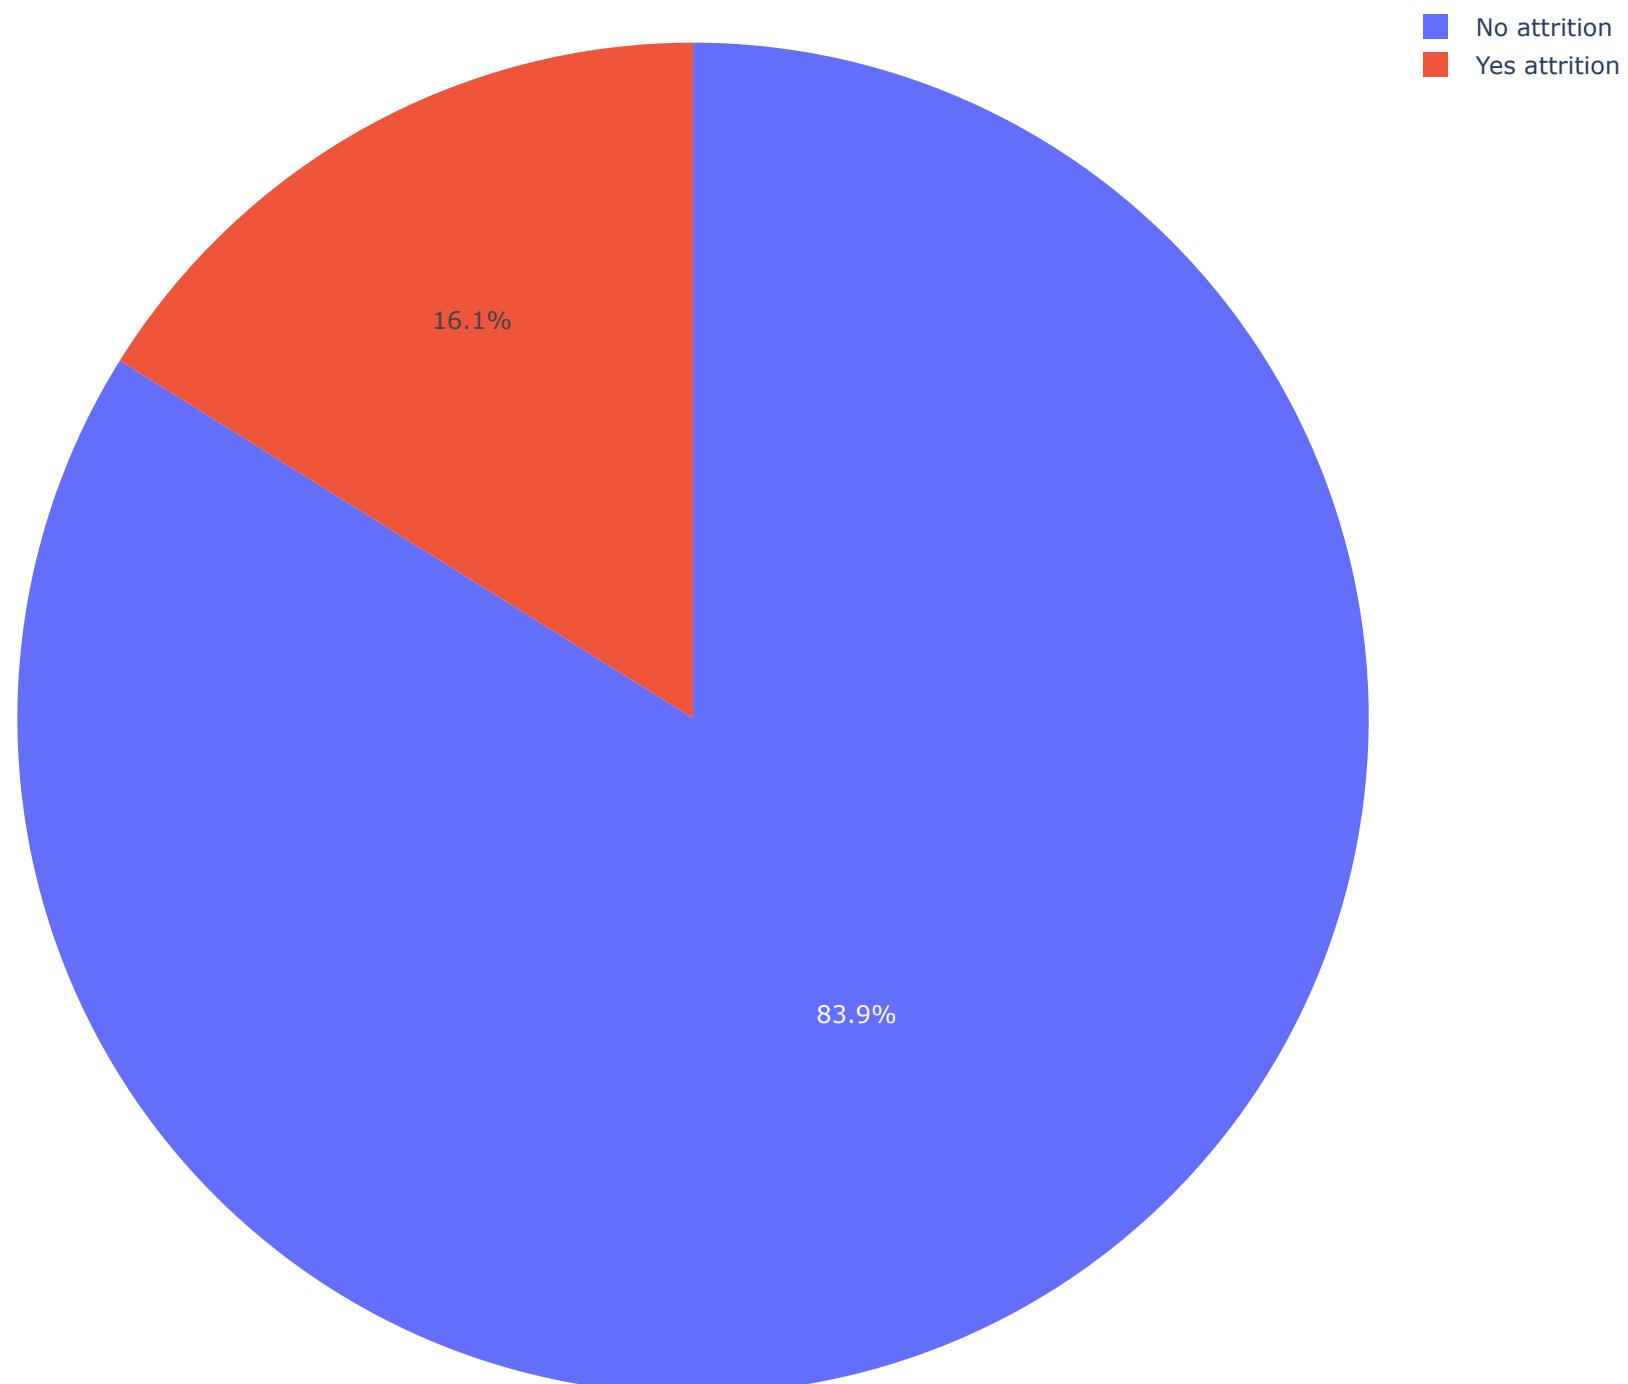
5 rows × 52 columns

```
In [28]: #Splitting each attrition class for more details
nAttrition = data[(data["Attrition"] == 0)]
yAttrition = data[(data["Attrition"] == 1)]
```

In [29]: #Look at the class distribution of attrition

```
trace = [go.Pie(  
    values = [len(nAttrition), len(yAttrition)],  
    labels=["No attrition", "Yes attrition"]  
)  
layout = go.Layout(  
    width = 900,  
    height = 900  
)  
fig = go.Figure(data = trace, layout = layout)  
fig.update_layout(title_text = "Attrition class distribution")  
fig.show()  
#plt.plot(fig)
```

Attrition class distribution



There's a large class imbalance so if an accurate model is to be trained from this dataset, resampling will be required

```
In [30]: z = data.describe()
display(z)
```

	Attrition	Age	DailyRate	DistanceFromHome	Education	EnvironmentSatisfaction	HourlyRate	JobInvolvement	JobLevel	JobSatisf
count	1470.000000	1470.000000	1470.000000	1470.000000	1470.000000	1470.000000	1470.000000	1470.000000	1470.000000	1470.000000
mean	0.161224	36.923810	802.485714	9.192517	2.912925	2.721769	65.891156	2.729932	2.063946	2.7
std	0.367863	9.135373	403.509100	8.106864	1.024165	1.093082	20.329428	0.711561	1.106940	1.1
min	0.000000	18.000000	102.000000	1.000000	1.000000	1.000000	30.000000	1.000000	1.000000	1.0
25%	0.000000	30.000000	465.000000	2.000000	2.000000	2.000000	48.000000	2.000000	1.000000	2.0
50%	0.000000	36.000000	802.000000	7.000000	3.000000	3.000000	66.000000	3.000000	2.000000	3.0
75%	0.000000	43.000000	1157.000000	14.000000	4.000000	4.000000	83.750000	3.000000	3.000000	4.0
max	1.000000	60.000000	1499.000000	29.000000	5.000000	4.000000	100.000000	4.000000	5.000000	4.0

8 rows × 52 columns

```
In [31]: yInfo = yAttrition.describe()
nInfo = nAttrition.describe()
yInfo = yInfo.drop(yInfo.index[0])
nInfo = nInfo.drop(nInfo.index[0])
```

```
In [32]: def distPlot(var):
    #Distribution plot
    dat1 = yAttrition[var]
    dat2 = nAttrition[var]
    histData = [dat1,dat2]
    groupLabel = ["Yes attrition", "No attrition"]

    autoBin = (z.loc["std", var])/8

    fig = ff.create_distplot(
        hist_data = histData,
        group_labels = groupLabel,
        show_rug=False,
        bin_size = autoBin
    )
    fig.update_layout(title_text = var)
    #plt.plot(fig)
    fig.show()
```

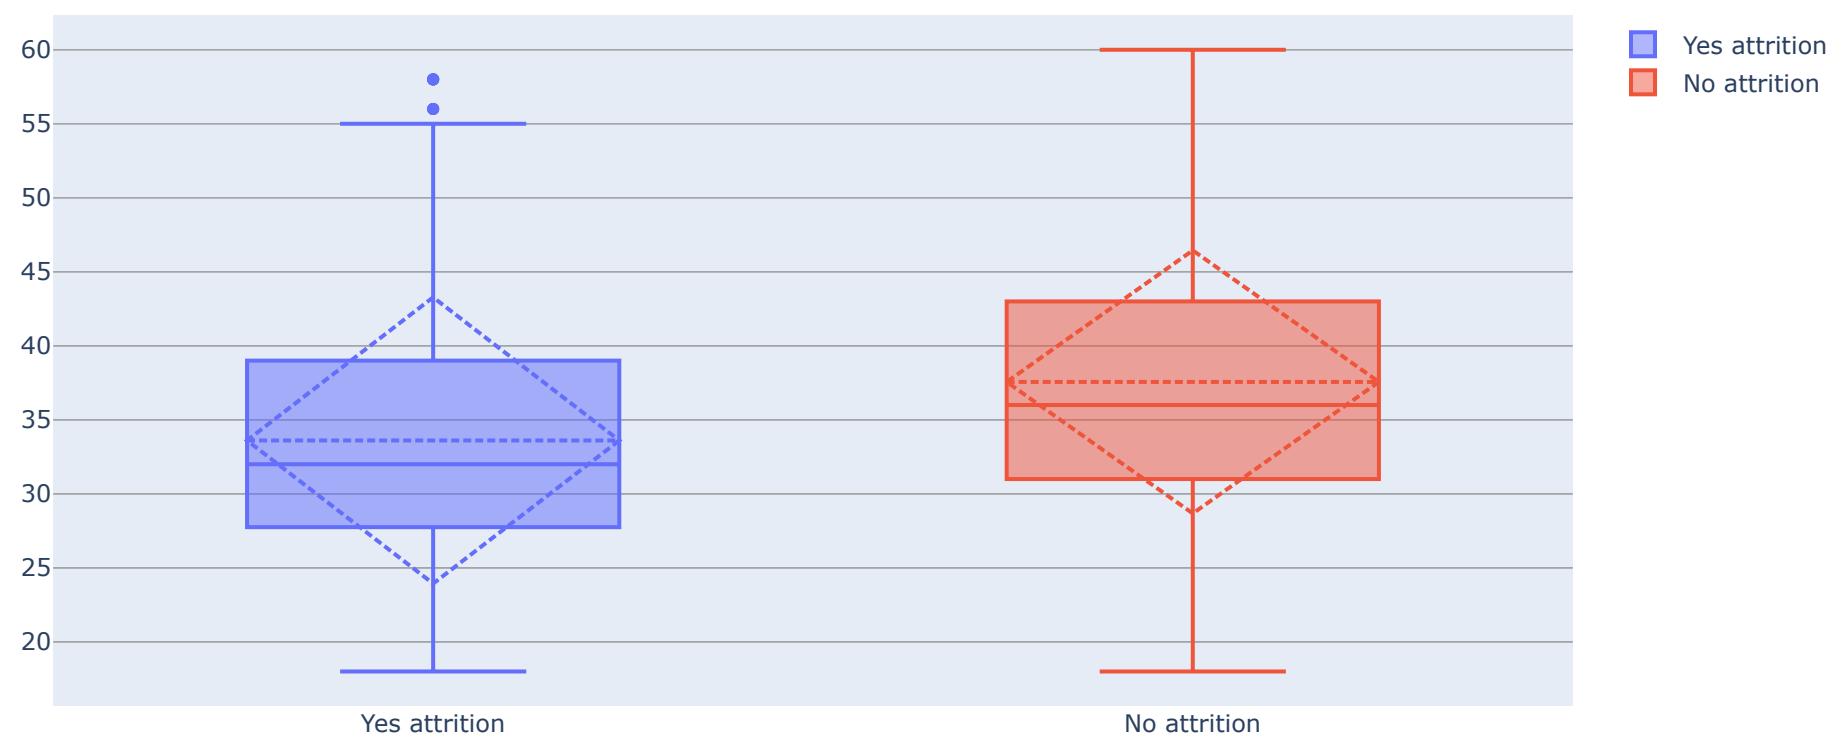
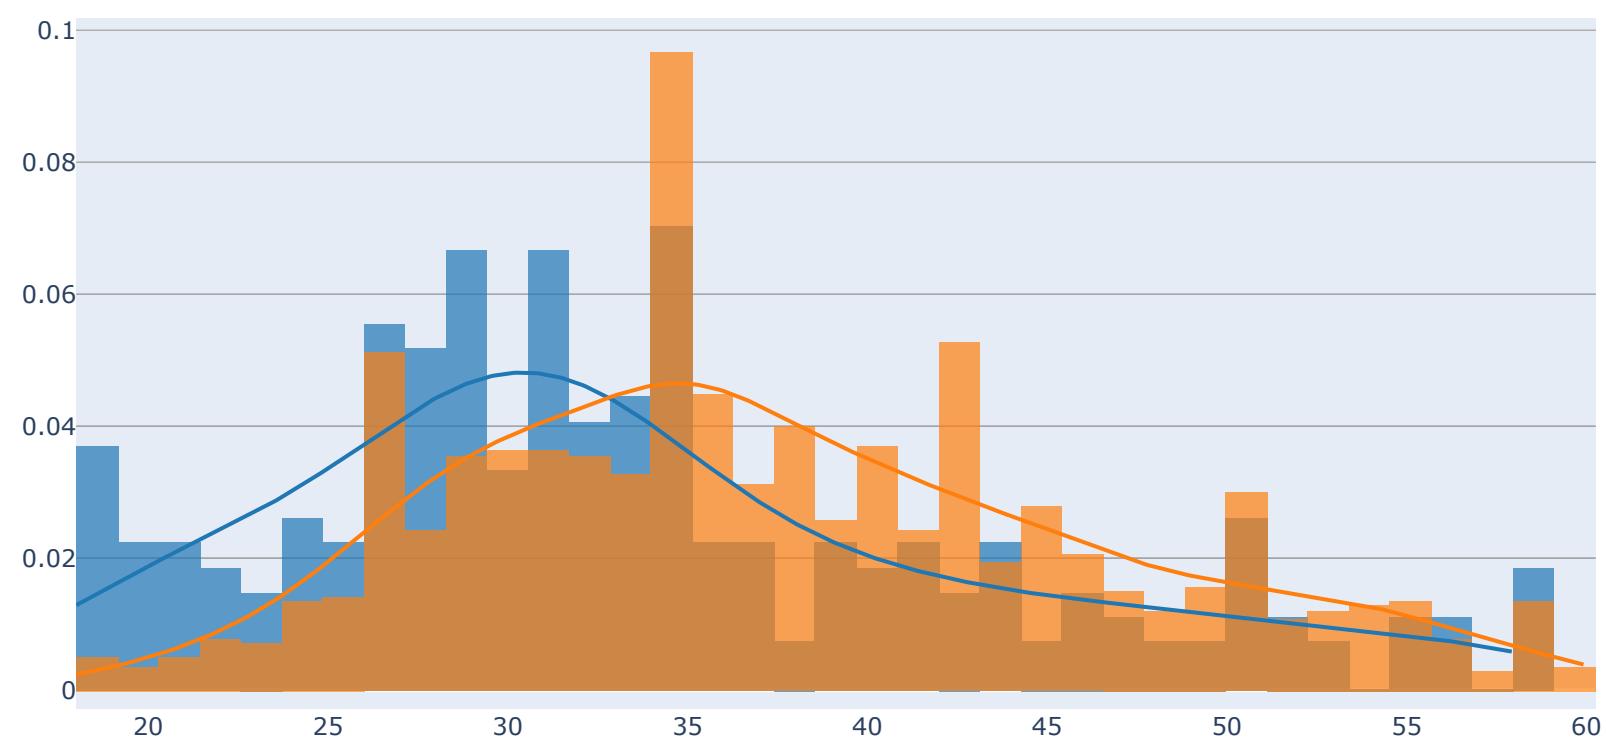
```
In [33]: def box(var):
    fig = go.Figure()
    yBox = yAttrition[var]
    nBox = nAttrition[var]
    fig.add_trace(go.Box(y=yBox, boxmean="sd", name="Yes attrition"))
    fig.add_trace(go.Box(y=nBox, boxmean="sd", name="No attrition"))
    #plt.plot(fig)
    fig.show()
```

```
In [34]: def comp(var):
    yMean = yInfo.loc["mean",var]
    nMean = nInfo.loc["mean",var]
    print("Yes-Attrition mean: \t" + str(yMean) + " ± " + str(yInfo.loc["std",var]))
    print("No-Attrition mean: \t" + str(nMean) + " ± " + str(nInfo.loc["std",var]))
    print("Absolute mean difference: \t" + str(abs(yMean-nMean)))
```

```
In [35]: def vis(var):
    #Visualize distribution plot, box plot, and mean±std
    distPlot(var)
    box(var)
    comp(var)
```

In [36]: vis("Age")

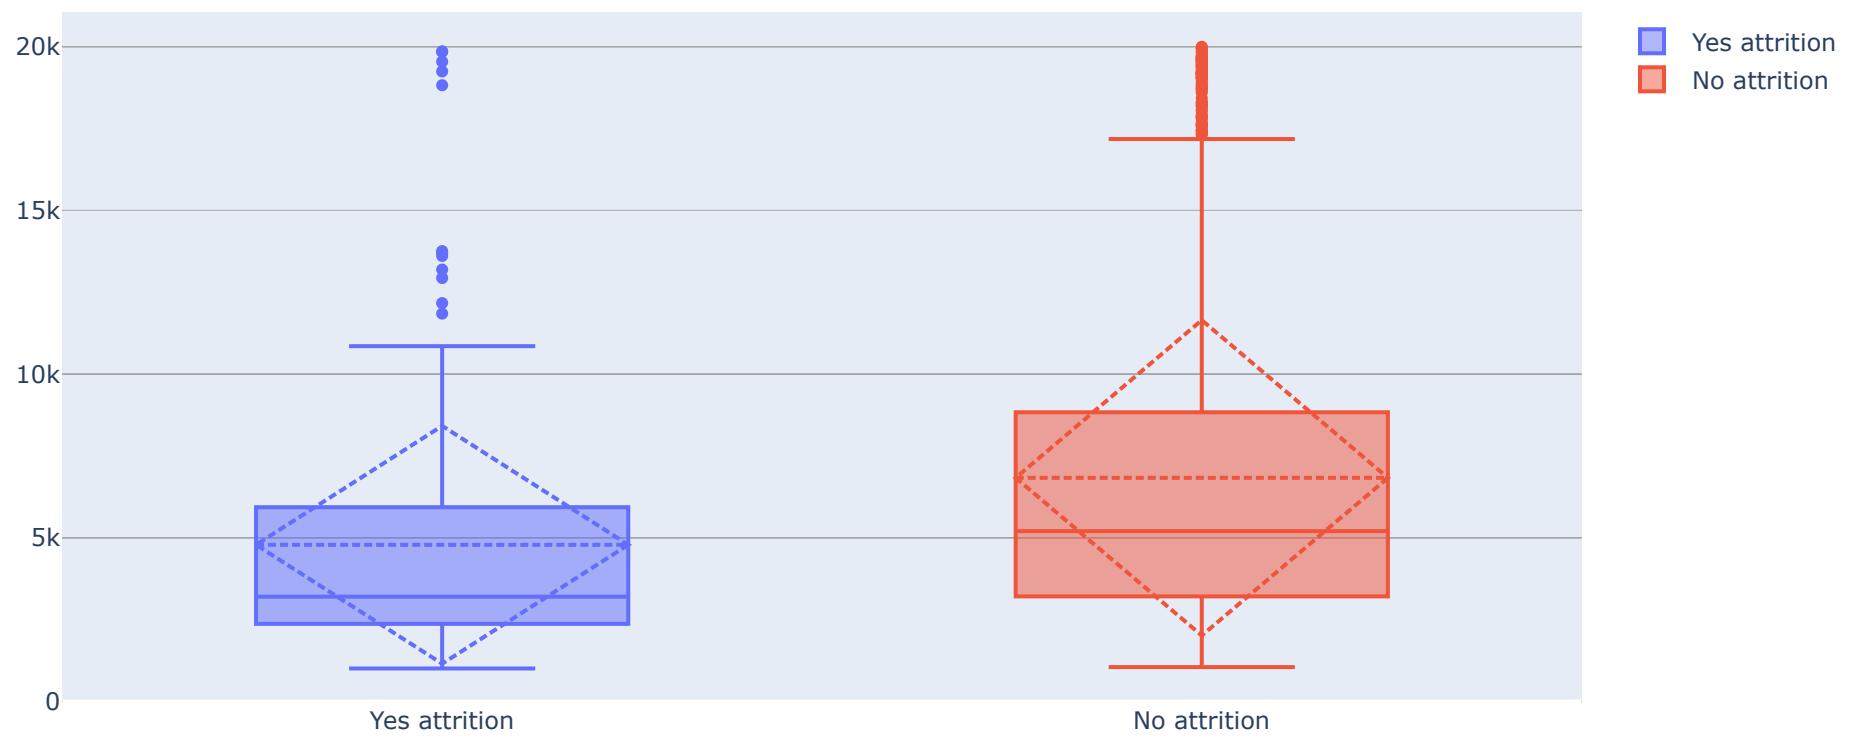
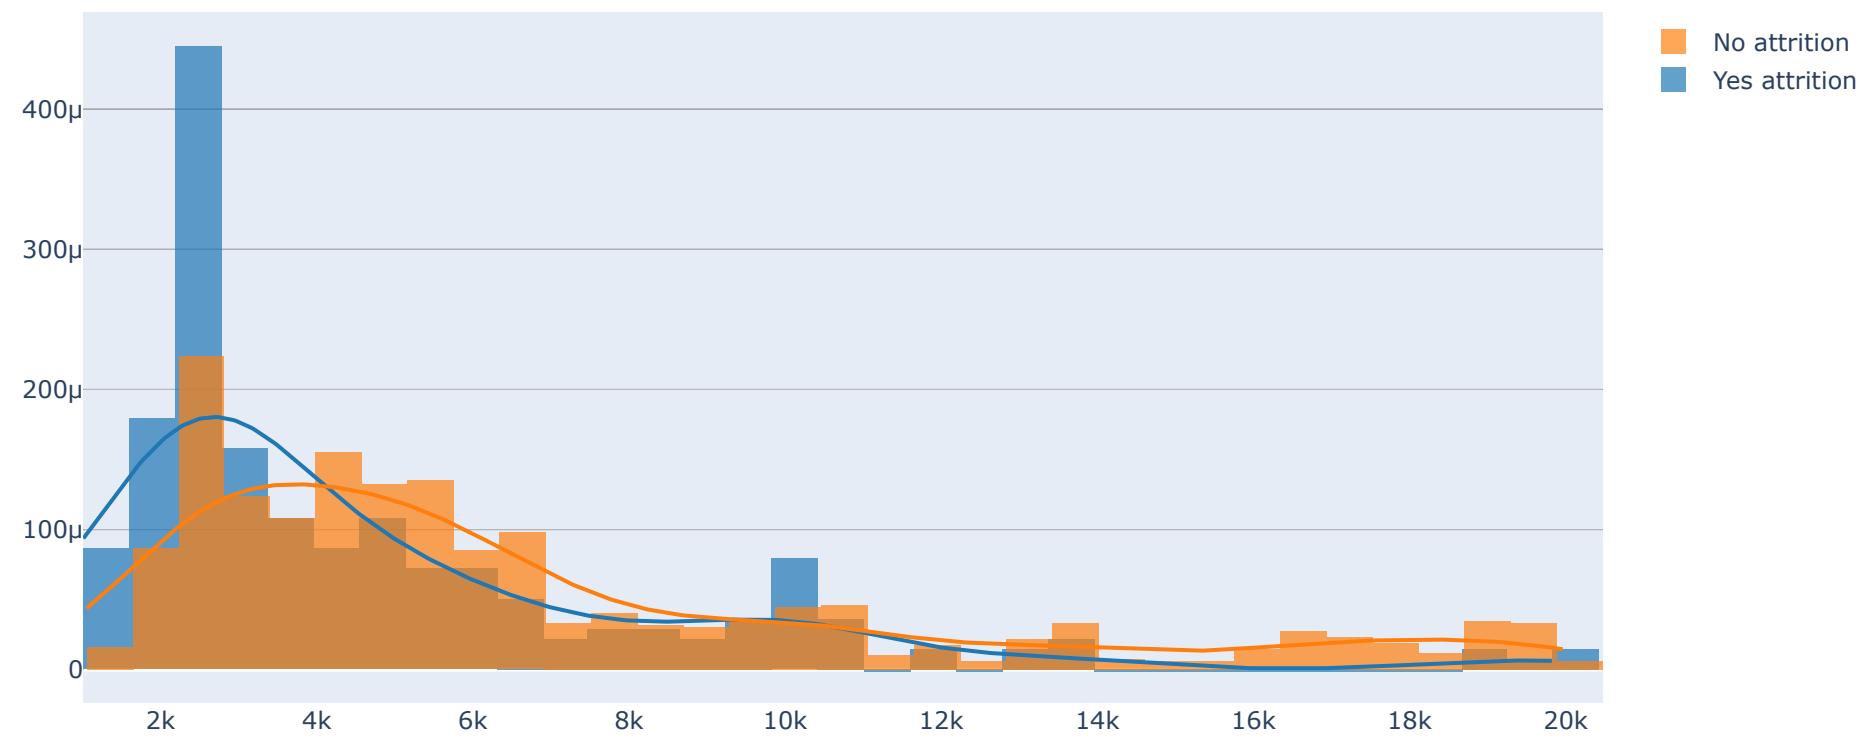
Age



Yes-Attrition mean: 33.607594936708864 ± 9.689349895351622
No-Attrition mean: 37.561232765612324 ± 8.888360024976546
Absolute mean difference: 3.95363782890346

In [37]: vis("MonthlyIncome")

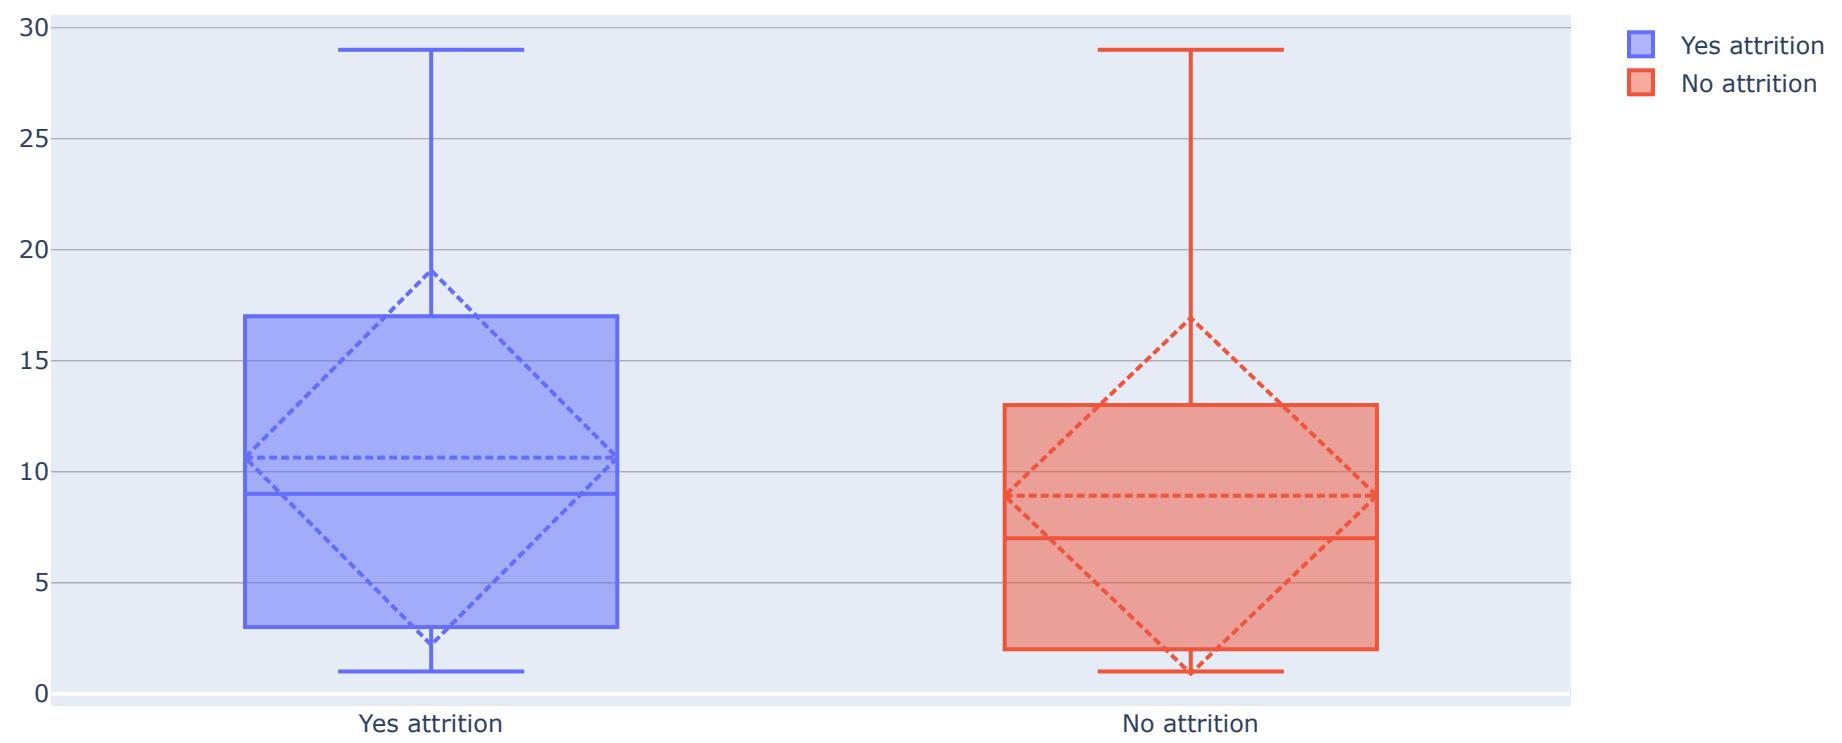
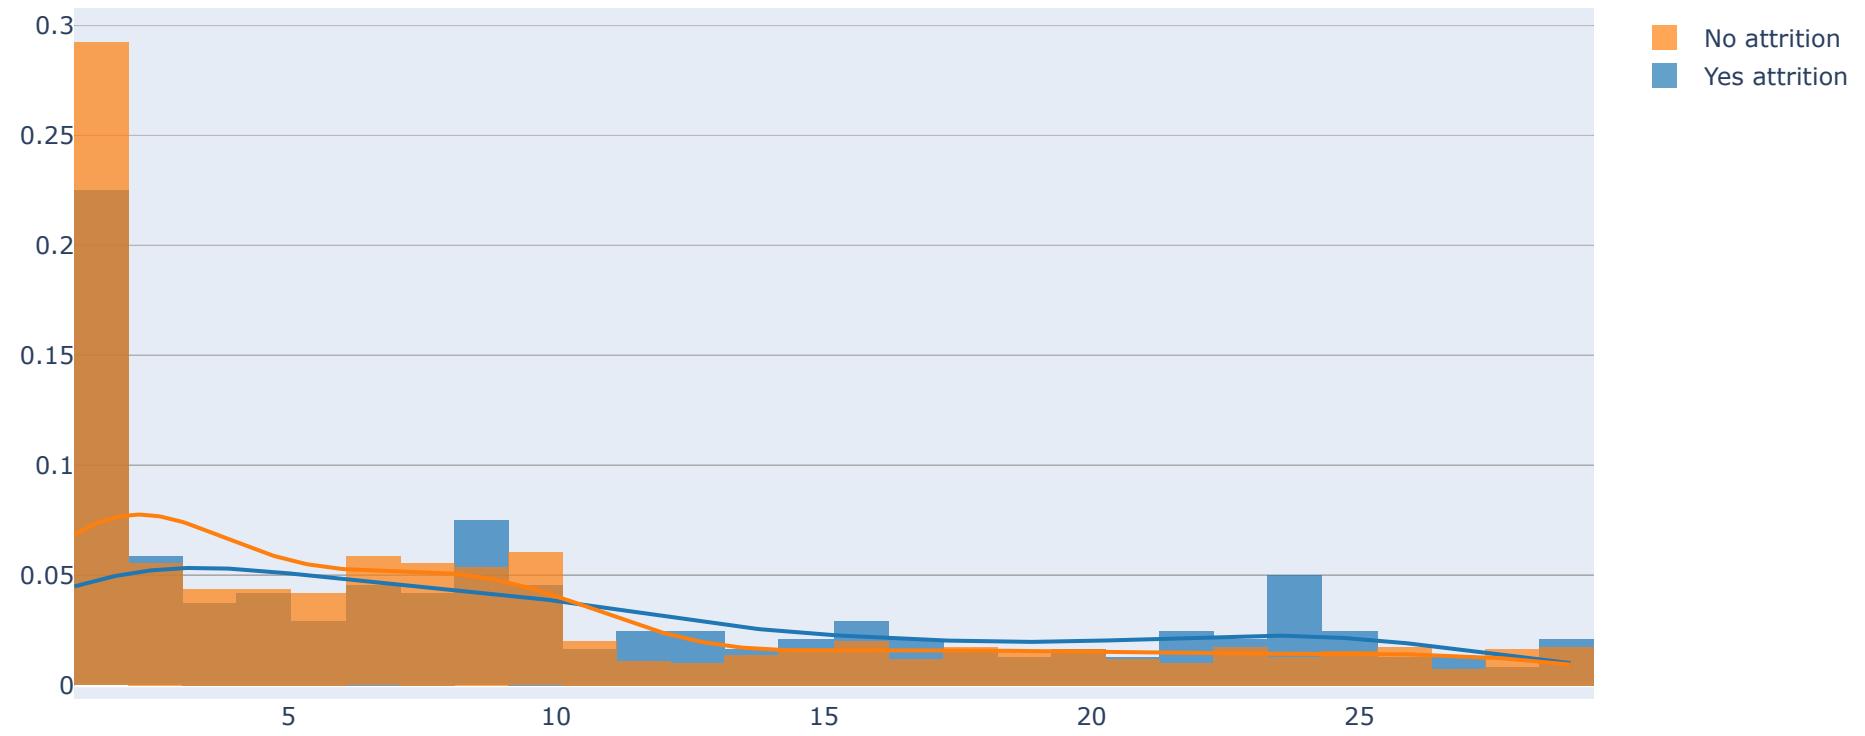
MonthlyIncome



Yes-Attrition mean: 4787.0928270042195 ± 3640.2103671038512
No-Attrition mean: 6832.739659367397 ± 4818.208000784485
Absolute mean difference: 2045.646832363177

In [38]: `vis("DistanceFromHome")`

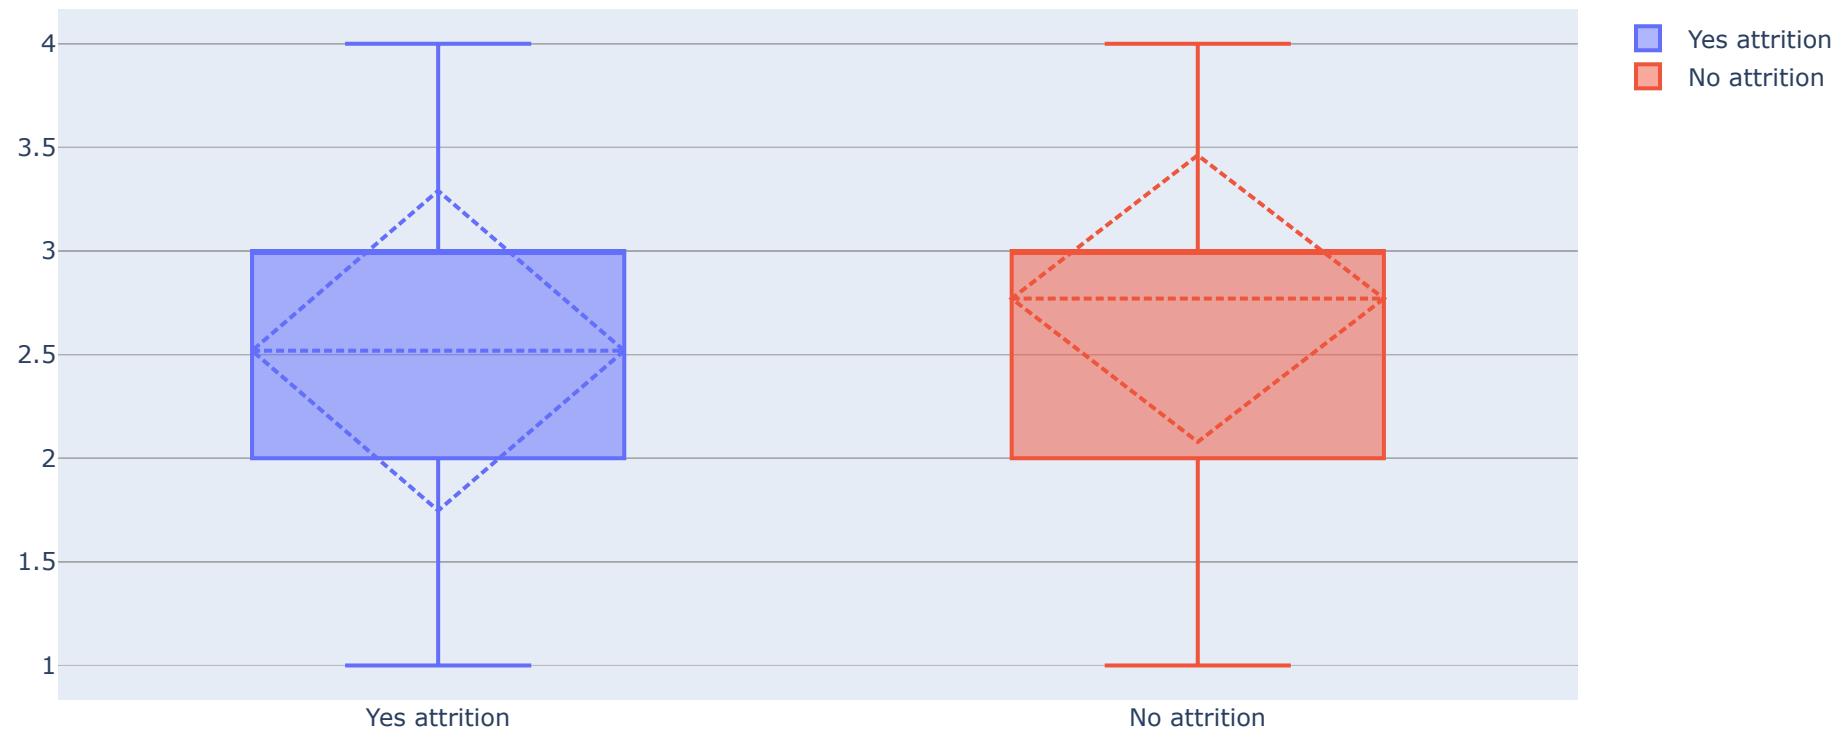
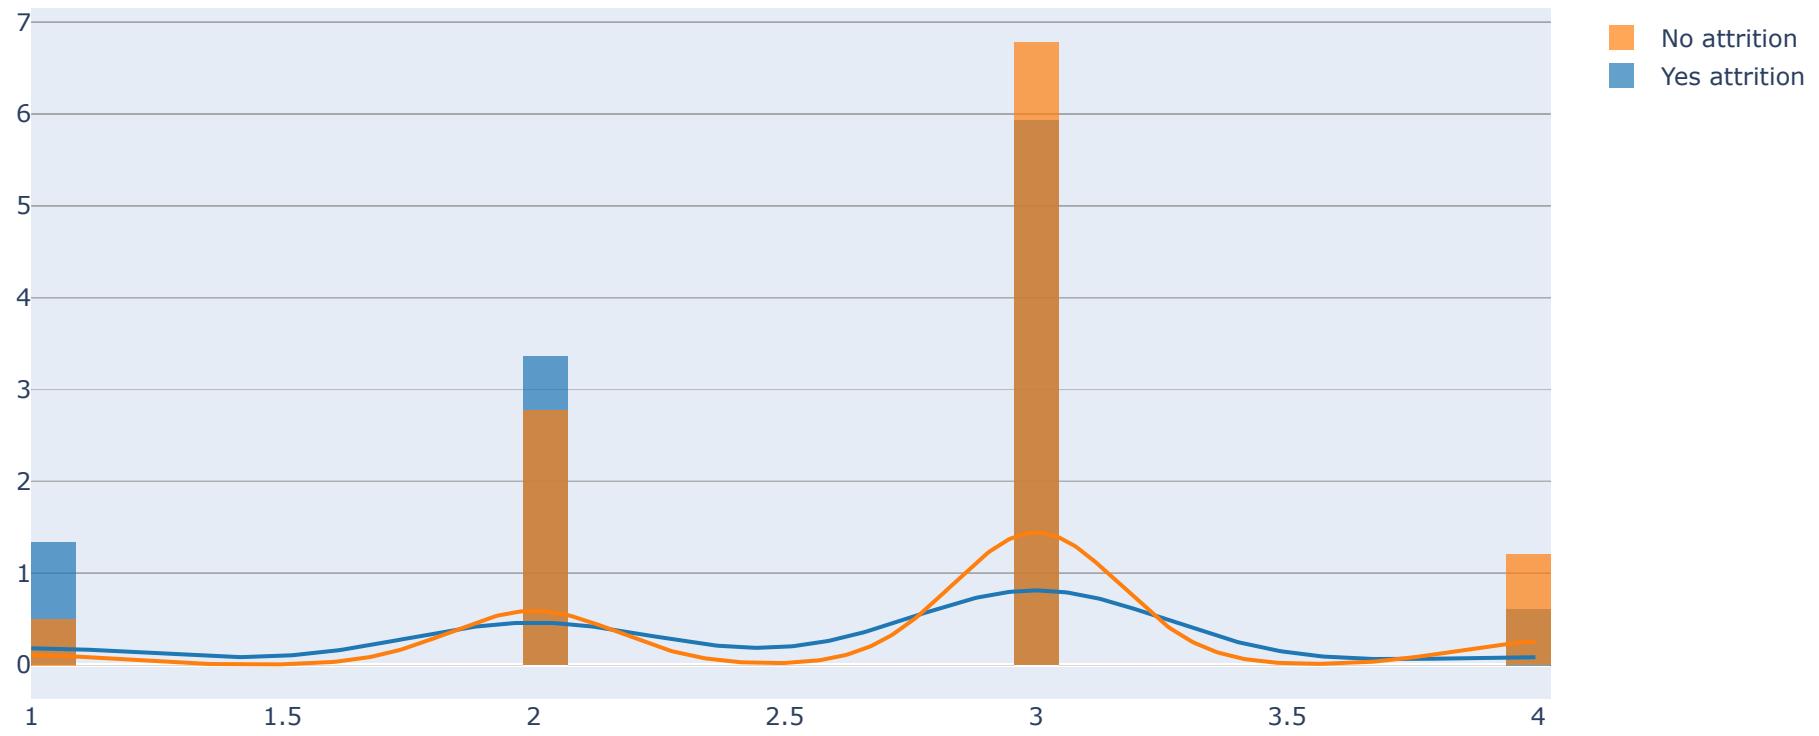
DistanceFromHome



Yes-Attrition mean: $10.632911392405063 \pm 8.452525269825024$
No-Attrition mean: $8.915652879156529 \pm 8.0126334854975$
Absolute mean difference: 1.7172585132485345

In [39]: vis("JobInvolvement")

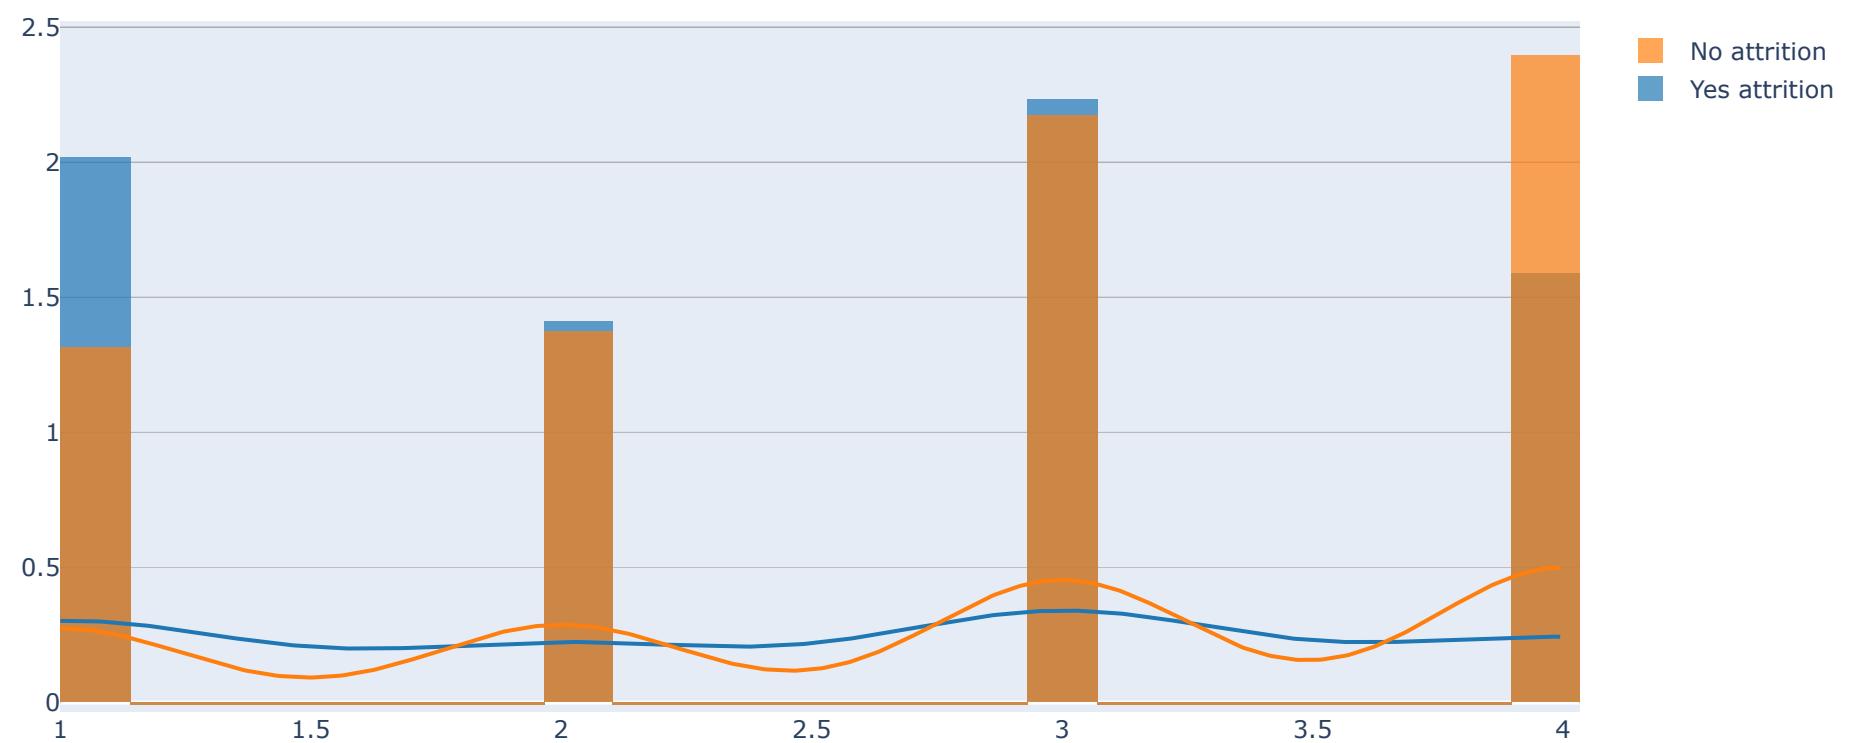
JobInvolvement



Yes-Attrition mean: 2.518987341772152 ± 0.7734047468056179
No-Attrition mean: 2.770478507704785 ± 0.6920498312234991
Absolute mean difference: 0.2514911659326331

In [40]: `vis("JobSatisfaction")`

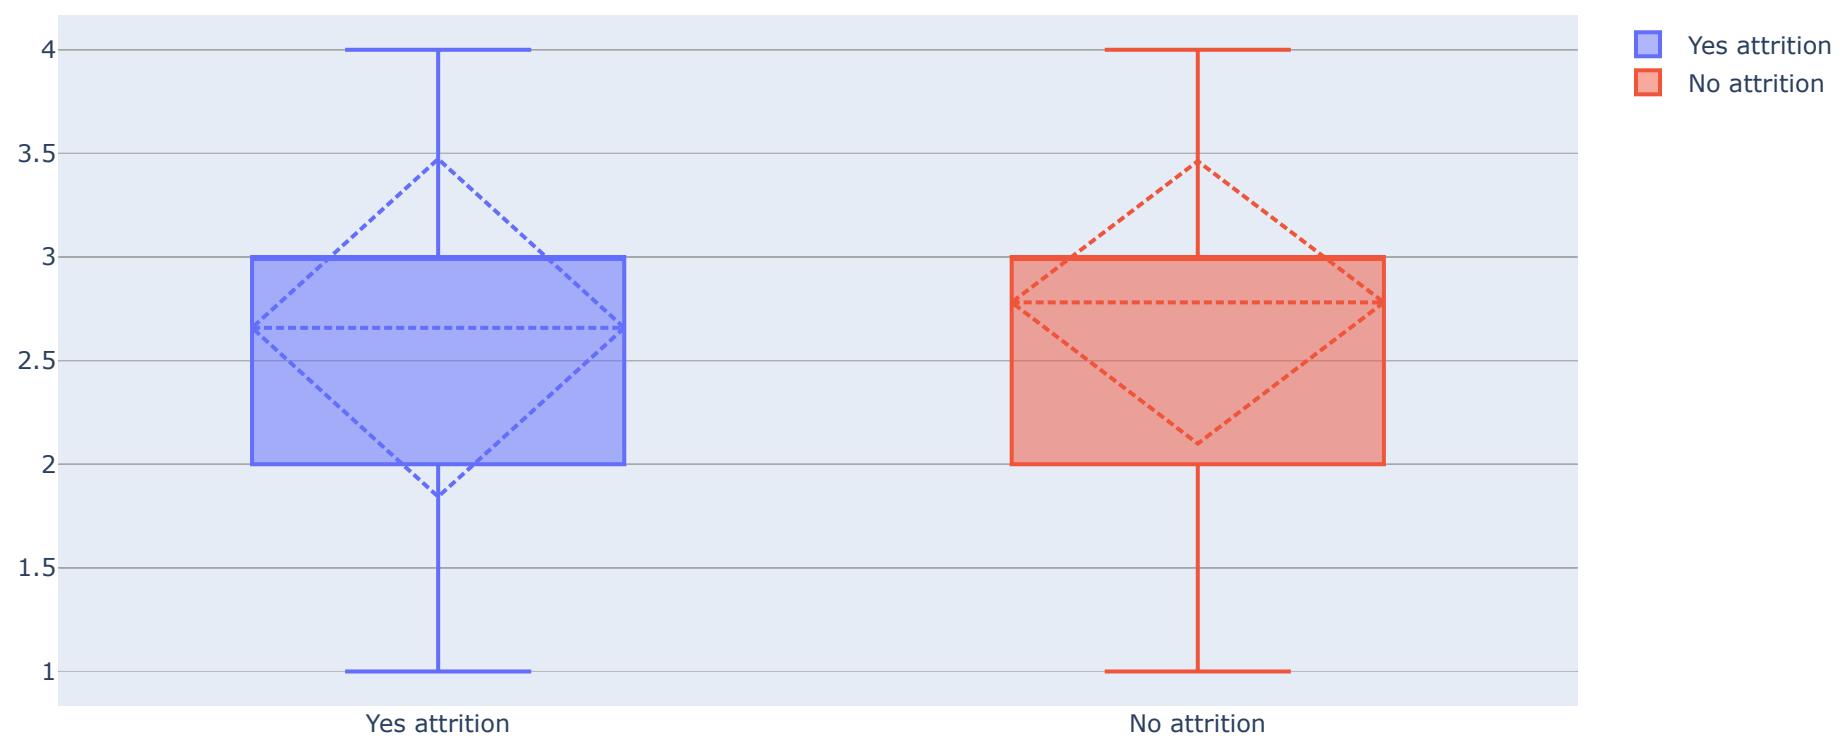
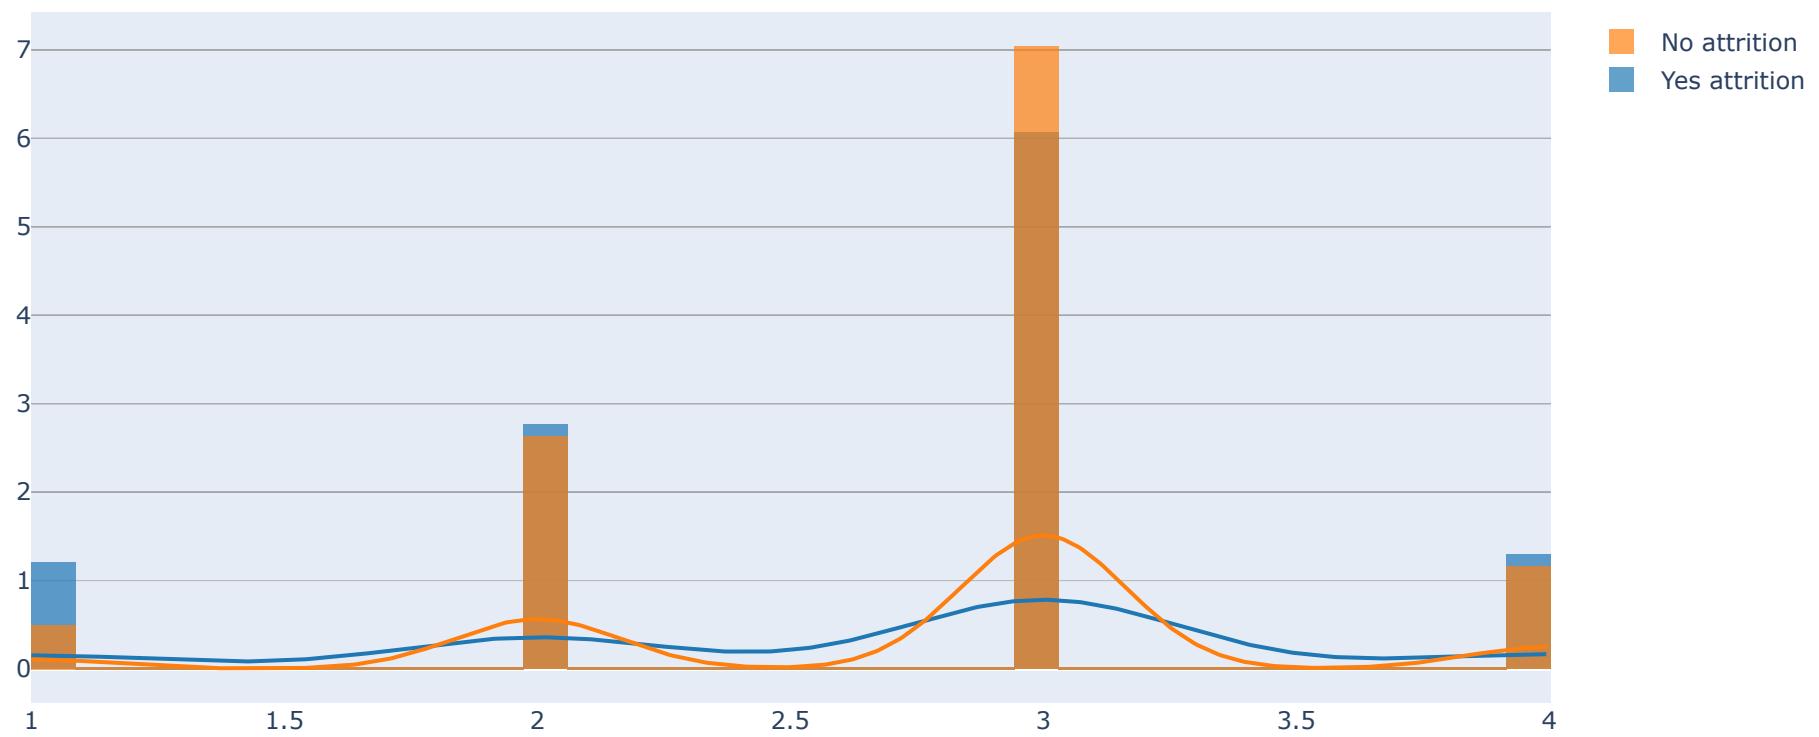
JobSatisfaction



Yes-Attrition mean: 2.4683544303797467 ± 1.11805797549084
No-Attrition mean: 2.778588807785888 ± 1.093277400503898
Absolute mean difference: 0.31023437740614135

In [41]: `vis("WorkLifeBalance")`

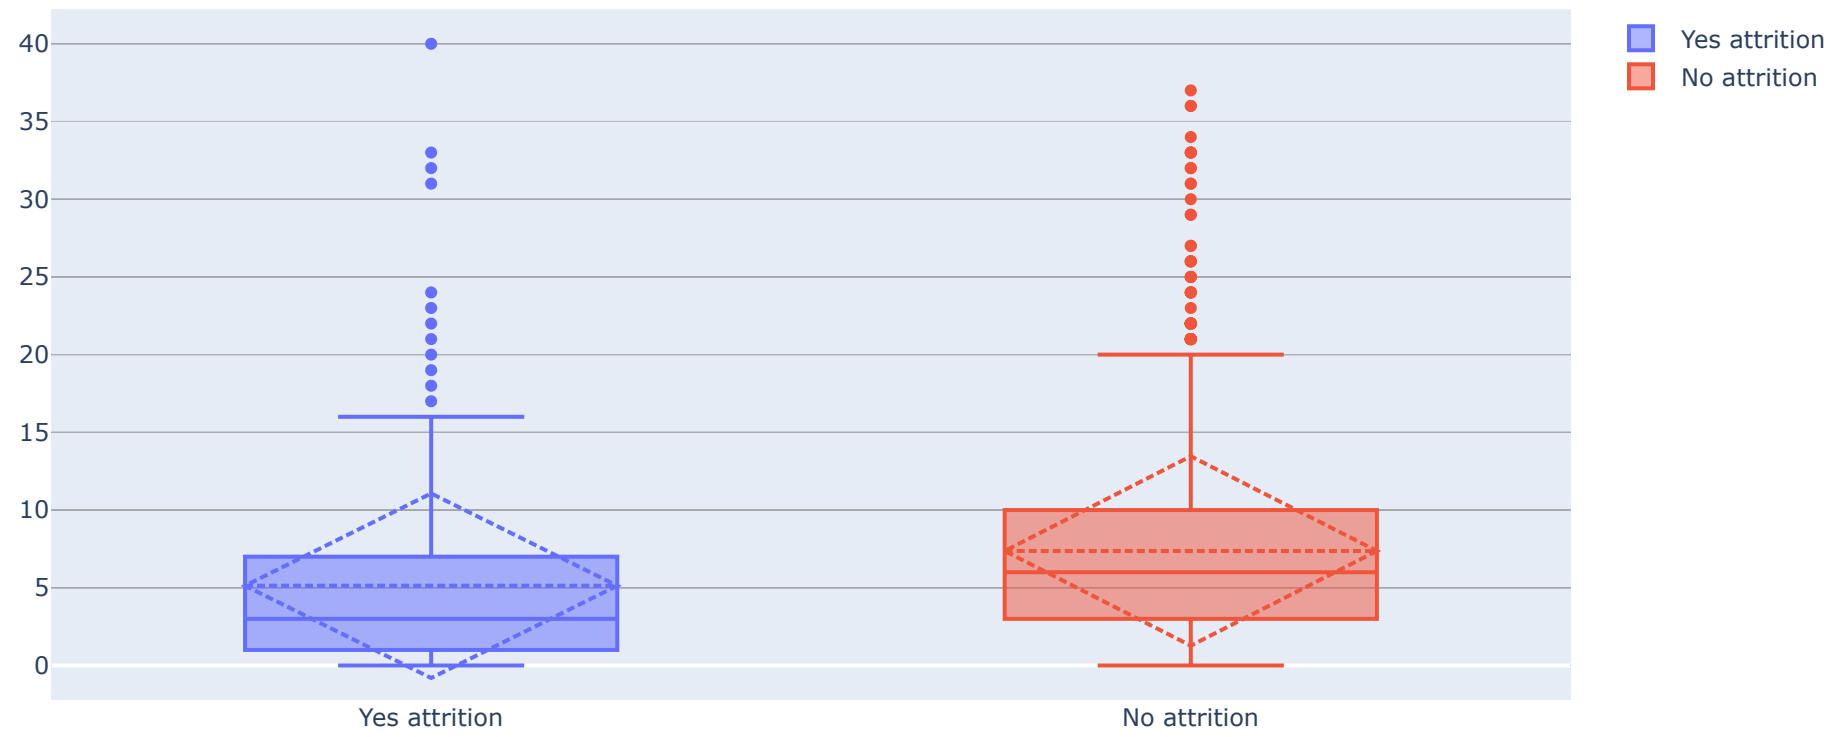
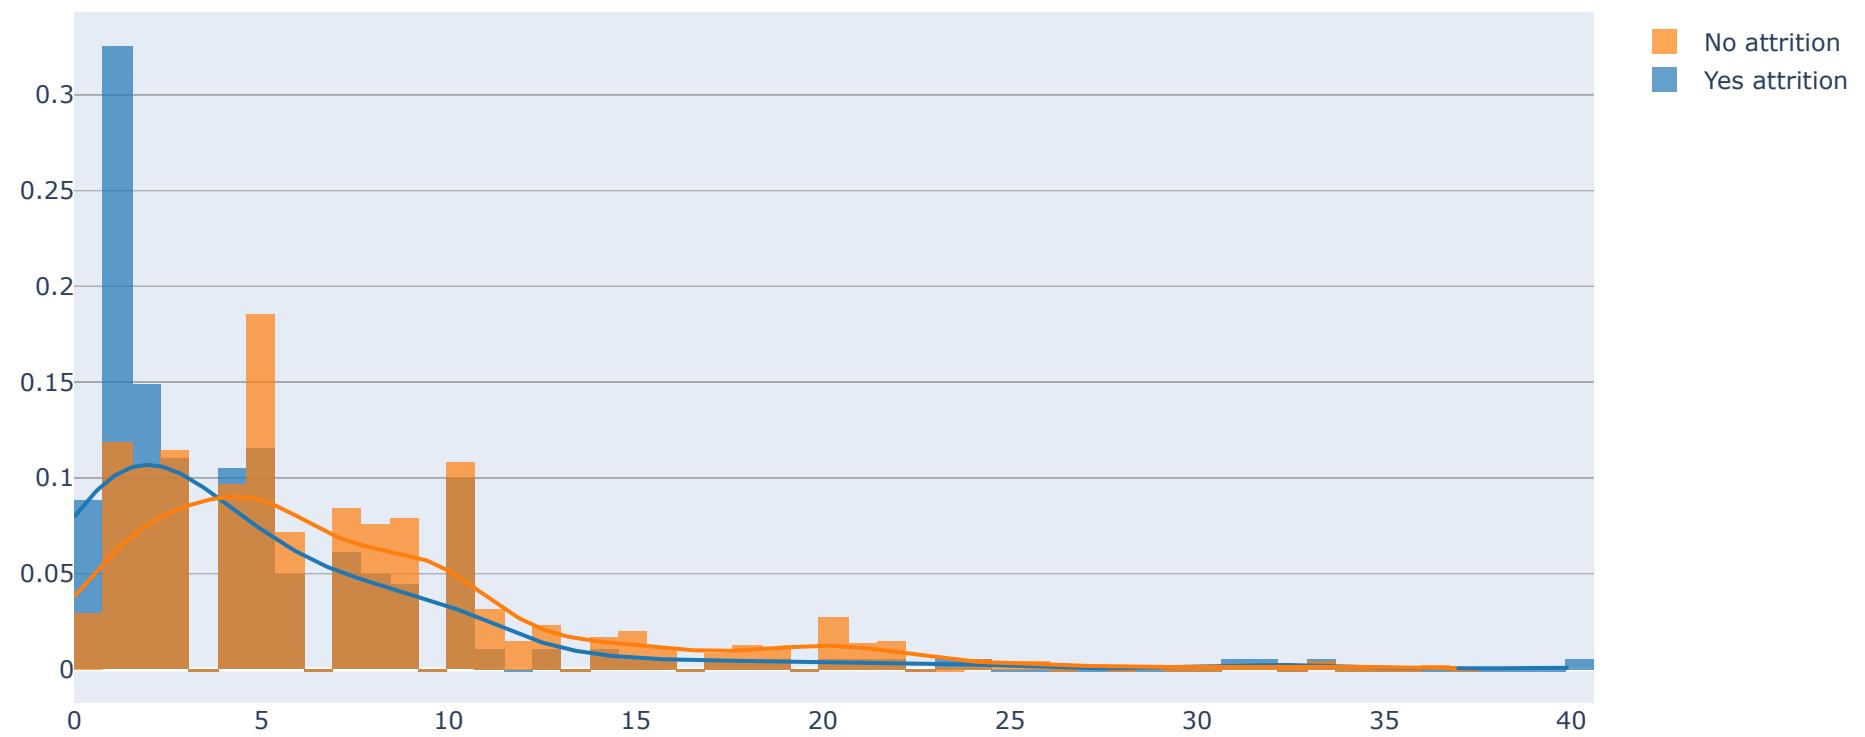
WorkLifeBalance



Yes-Attrition mean: 2.6582278481012658 ± 0.8164527856864083
No-Attrition mean: 2.781021897810219 ± 0.6819066520792771
Absolute mean difference: 0.12279404970895325

In [42]: vis("YearsAtCompany")

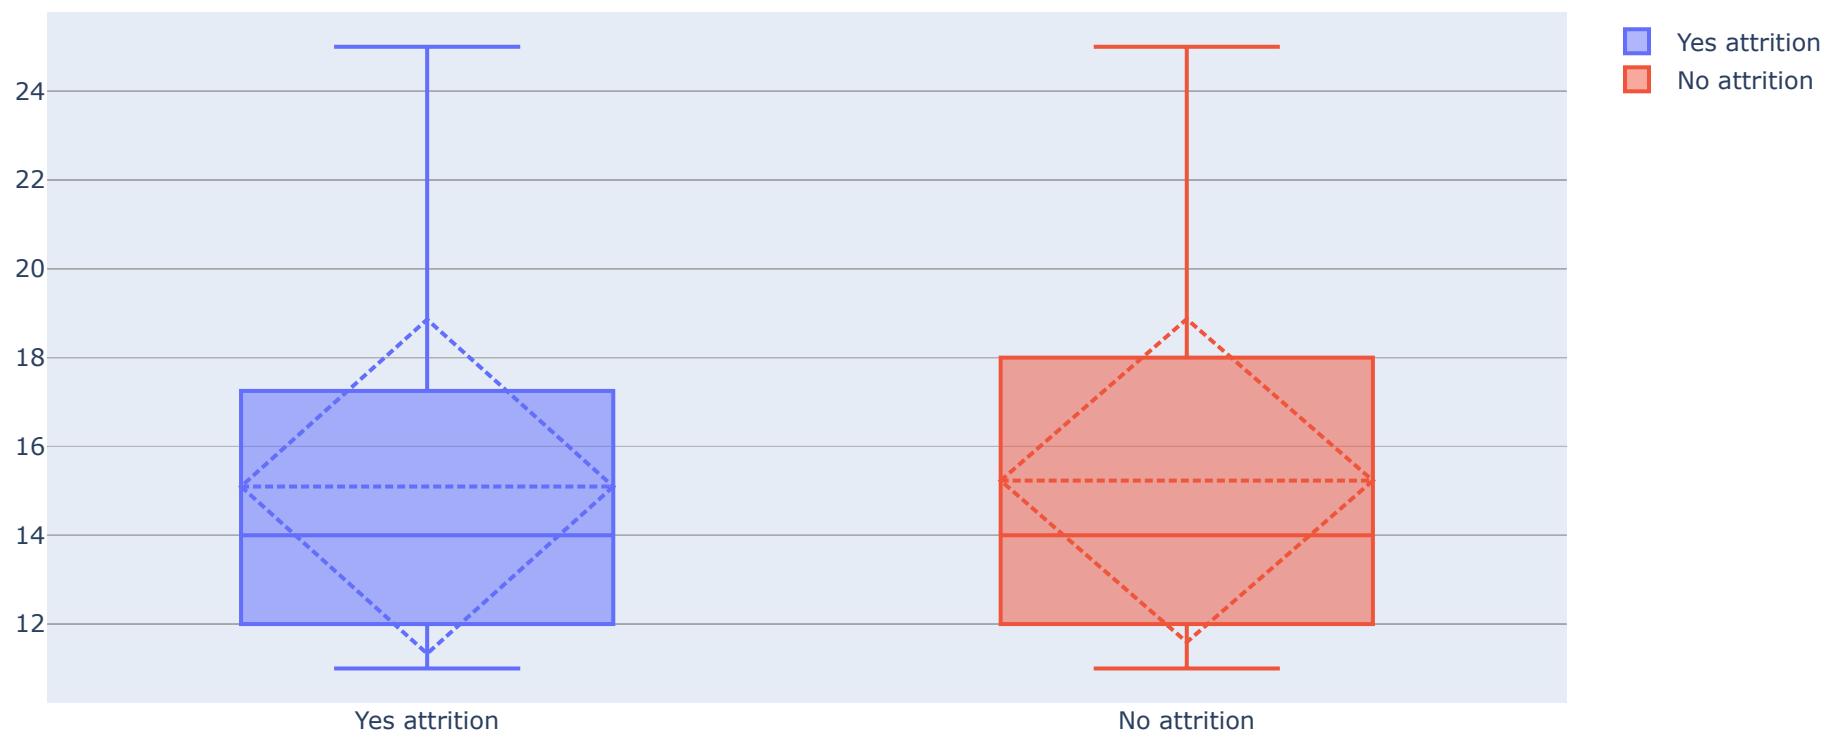
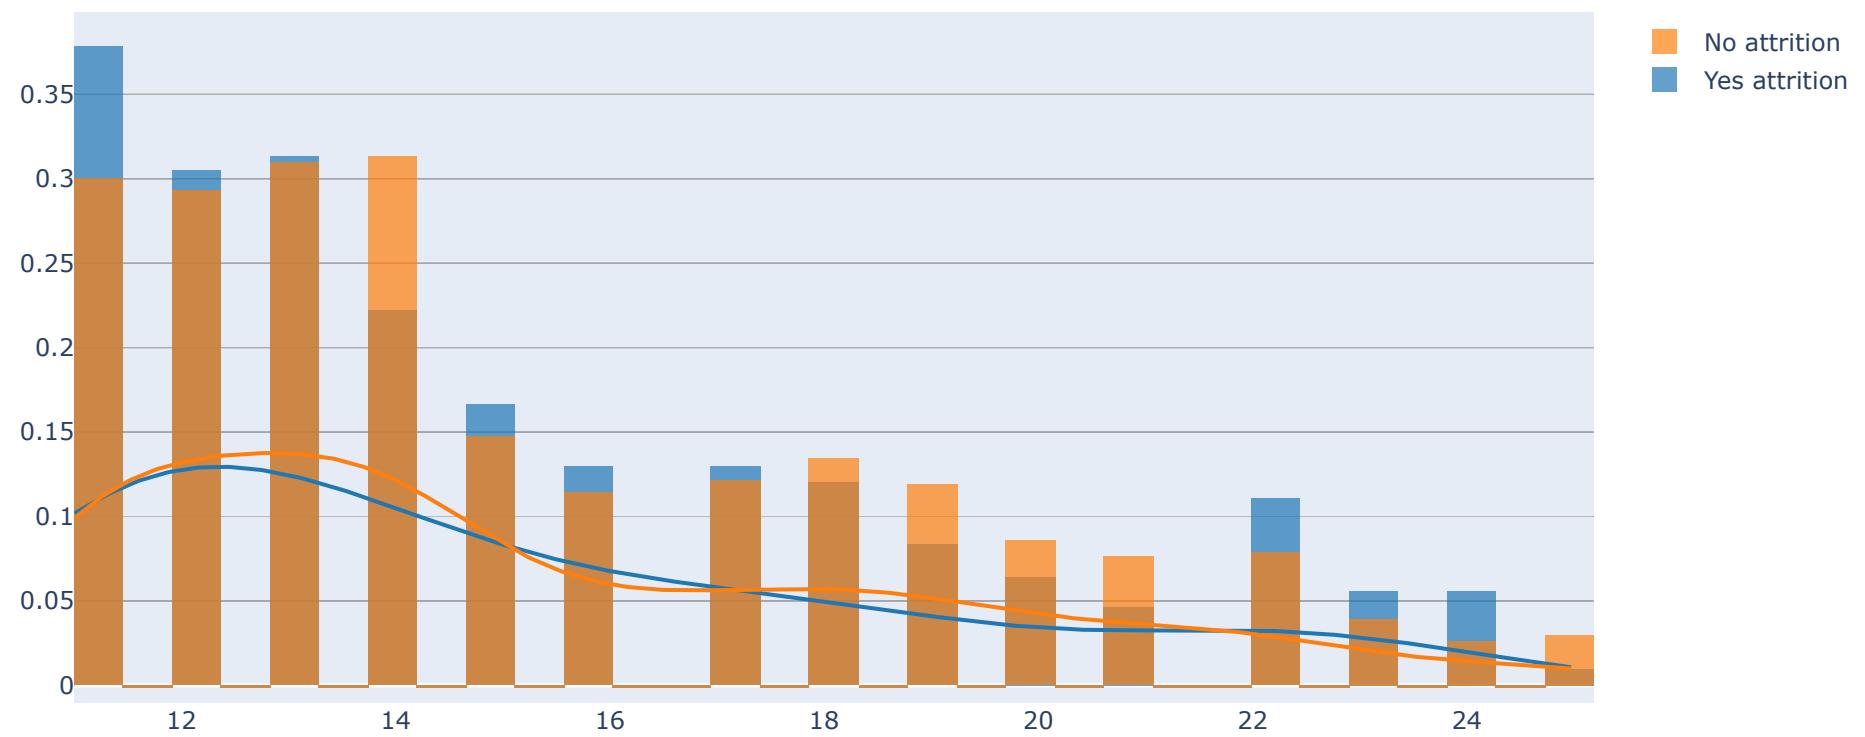
YearsAtCompany



Yes-Attrition mean: 5.1308016877637135 ± 5.949984029204934
No-Attrition mean: 7.369018653690187 ± 6.096298144398664
Absolute mean difference: 2.2382169659264735

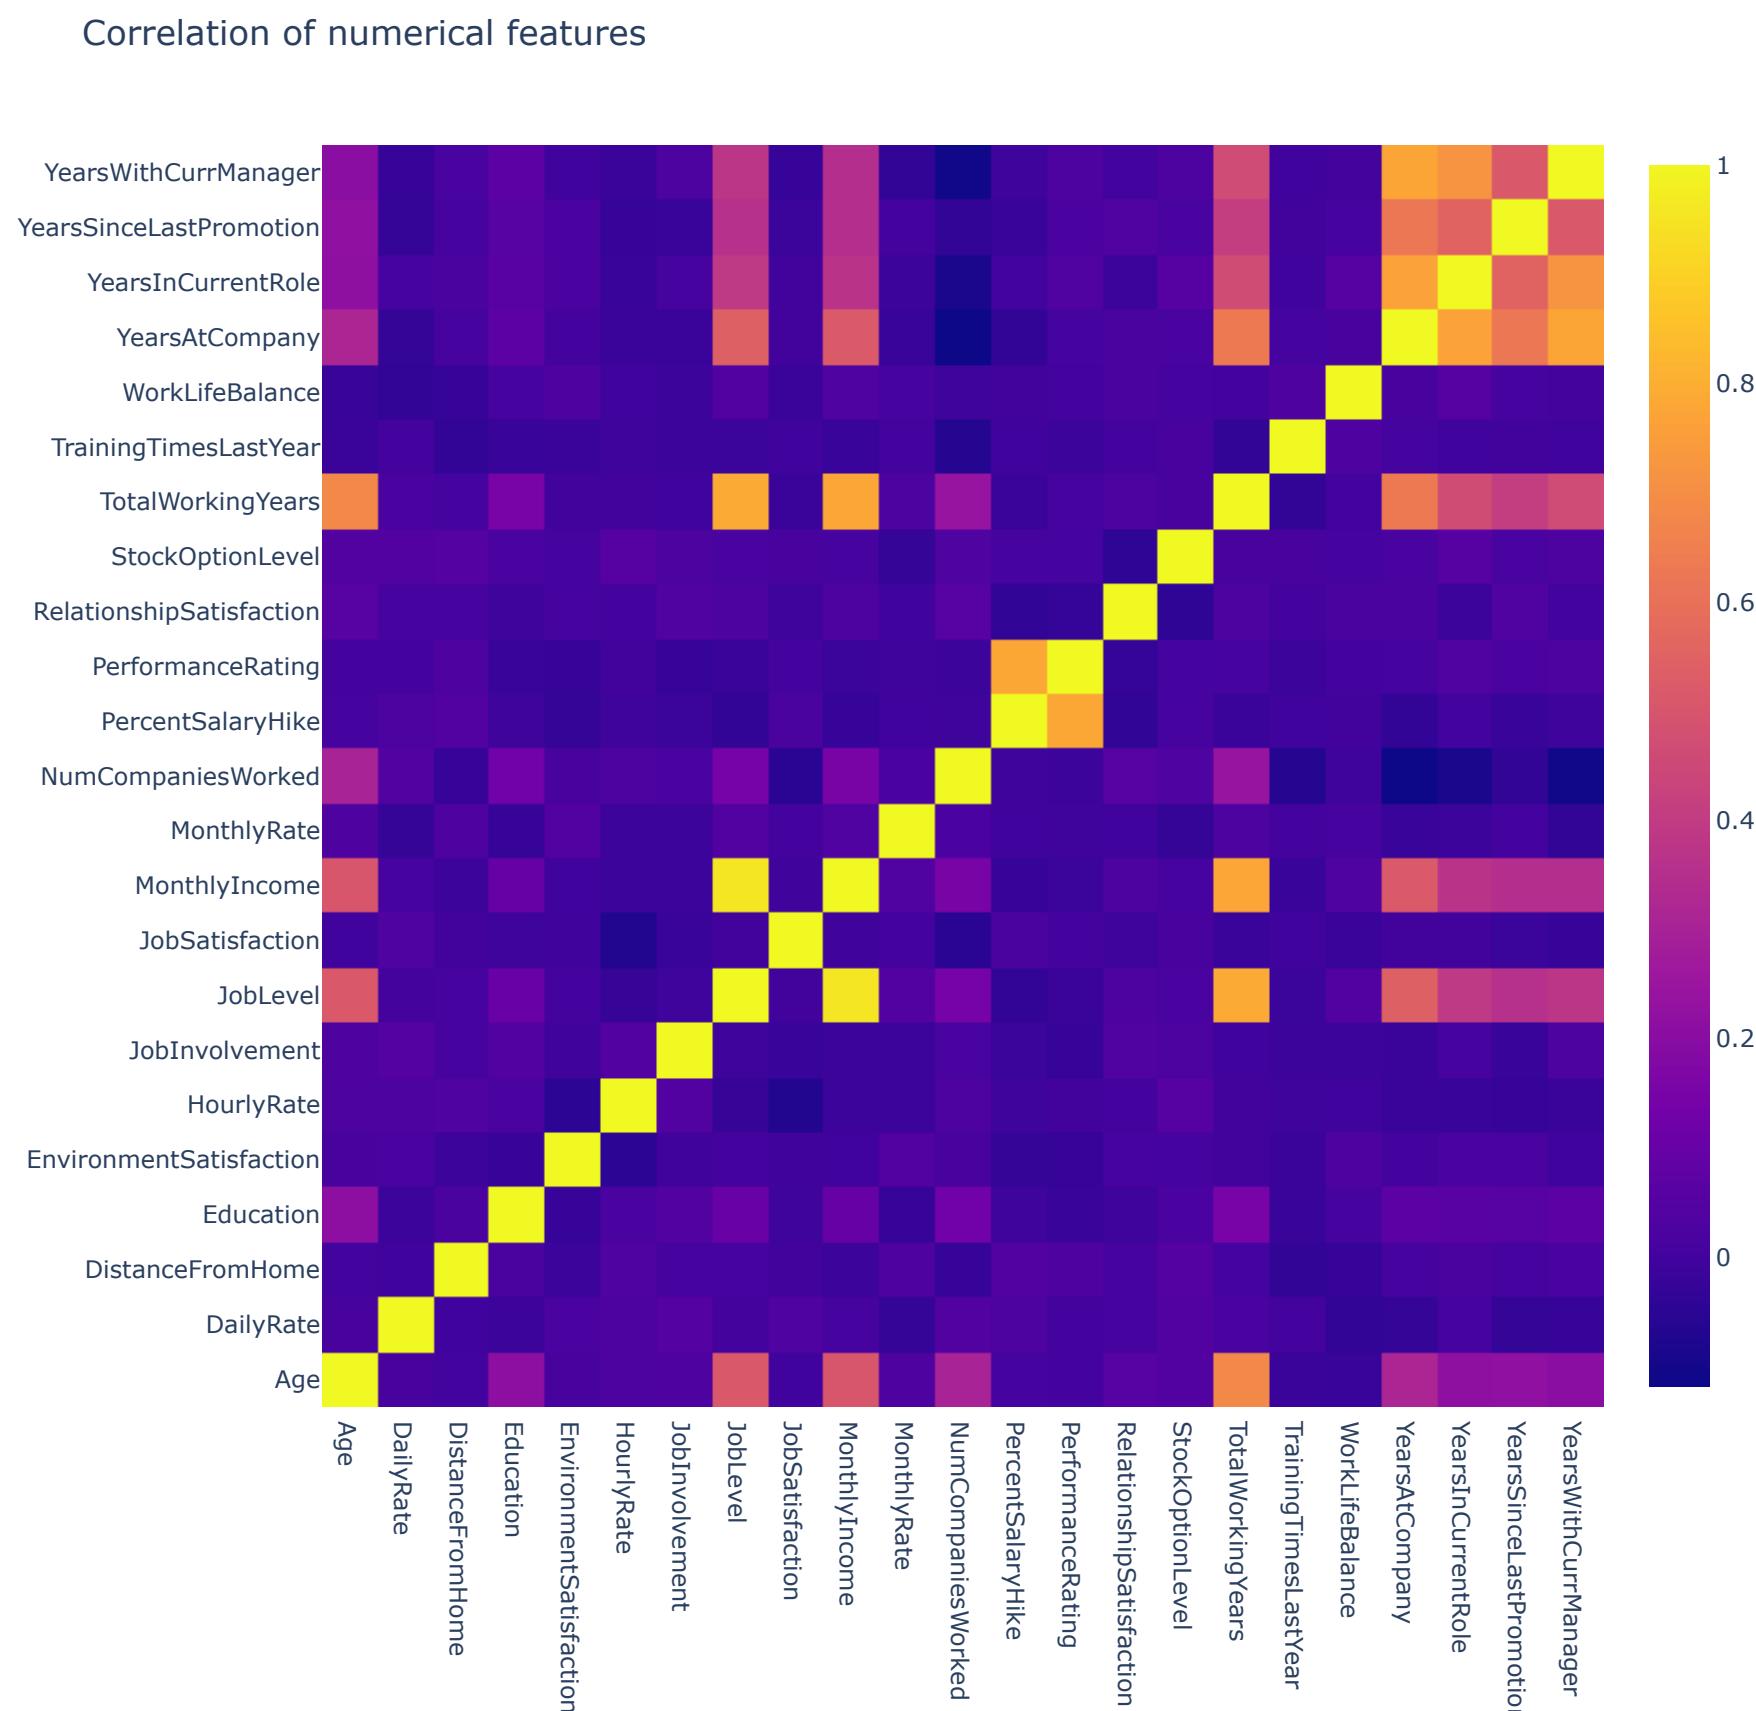
In [43]: vis("PercentSalaryHike")

PercentSalaryHike



Yes-Attrition mean: 15.09704641350211 ± 3.77029420045786
No-Attrition mean: 15.231143552311435 ± 3.639511269235044
Absolute mean difference: 0.1340971388093255

```
In [44]: trace = go.Heatmap(
    z = numFeatures.corr(),
    y = numFeatures.columns,
    x = numFeatures.columns
)
layout = go.Layout(
    title = "Correlation of numerical features",
    width = 900,
    height = 900
)
go.Figure(data = trace, layout = layout).show()
```



Some high correlation within features, possibly redundant

```
In [45]: #Threshold for correlation to be dropped ≥ 0.75
data = data.drop(columns=["TotalWorkingYears", "JobLevel", "PerformanceRating", "YearsInCurrentRole", "YearsWithCurrManager"])
```

```
In [46]: data.head()
```

Out[46]:

	Attrition	Age	DailyRate	DistanceFromHome	Education	EnvironmentSatisfaction	HourlyRate	JobInvolvement	JobSatisfaction	MonthlyIncome	...	JobRole
0	1	41	1102		1	2		2	94	3	4	5993 ...
1	0	49	279		8	1		3	61	2	2	5130 ...
2	1	37	1373		2	2		4	92	2	3	2090 ...
3	0	33	1392		3	4		4	56	3	3	2909 ...
4	0	27	591		2	1		1	40	3	2	3468 ...

5 rows × 47 columns

```
In [47]: #Splitting training and test set
data = data.drop(columns=["Attrition"])
xTrain, xTest, yTrain, yTest = train_test_split(data, label, test_size=0.30)
yTrain = np.ravel(yTrain)
yTest = np.ravel(yTest)
```

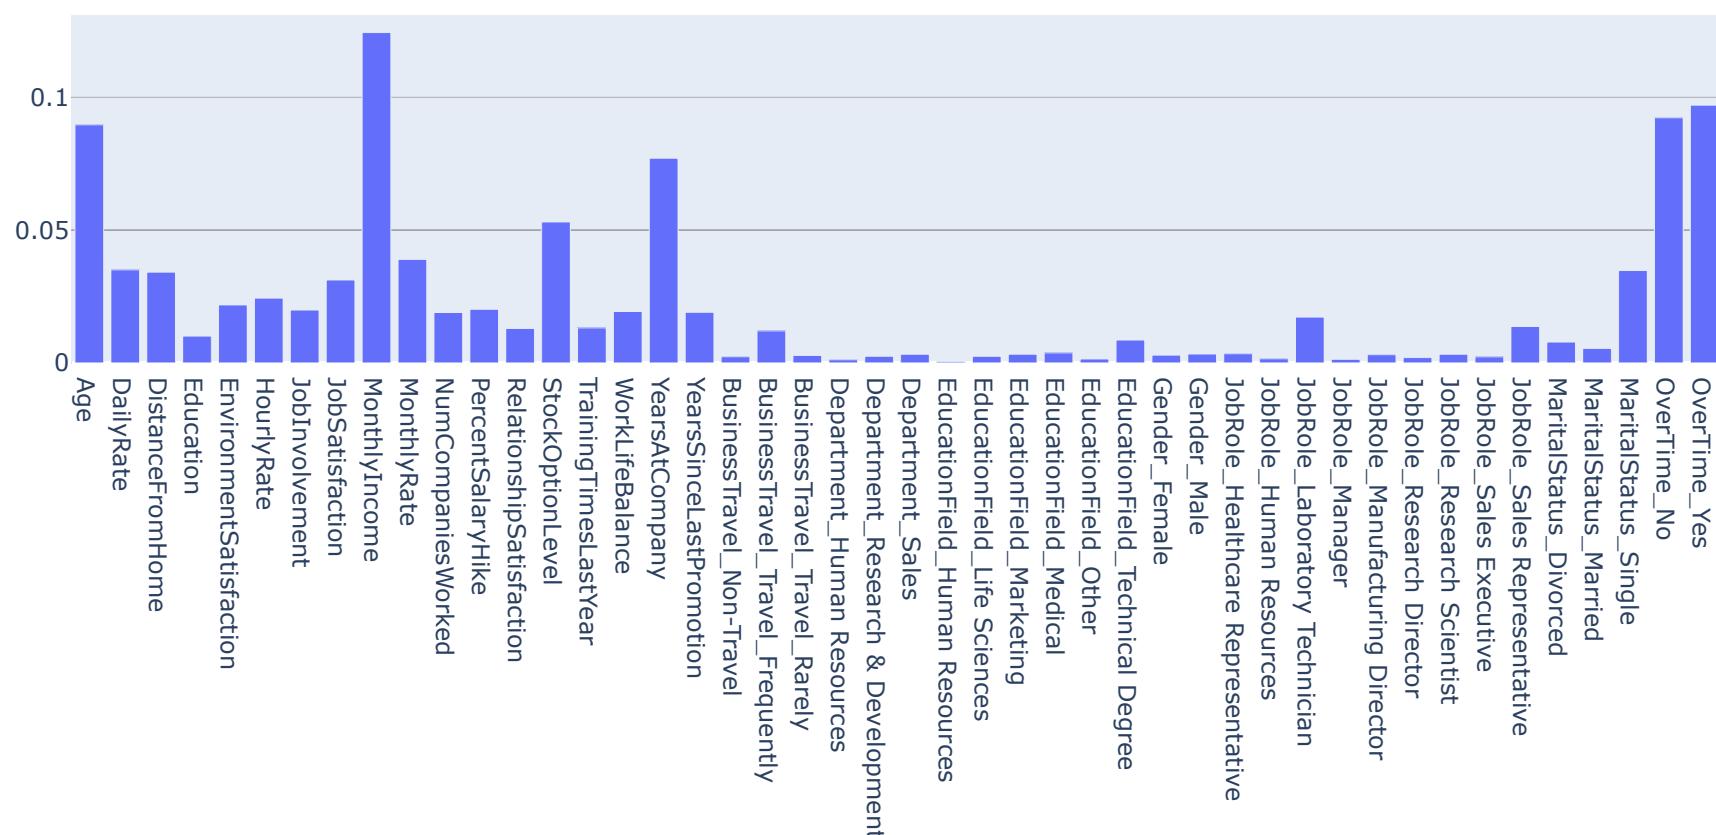
```
In [48]: def modelSummary(xTrain, yTrain, xTest, yTest):
    #Create a model
    rf = RandomForestClassifier(n_estimators=1000, max_depth=4)
    #Fit model to training set
    rf.fit(xTrain, yTrain)
    #Predict against test set
    predictions = rf.predict(xTest)
    #Evaluate model results
    print(confusion_matrix(yTest, predictions))
    print(classification_report(yTest,predictions, target_names=["No", "Yes"]))

    #Visualize the model feature importance ranks
    trace = go.Bar(
        y = rf.feature_importances_,
        x = data.columns.values
    )
    layout = go.Layout(
        title = "Feature importances"
    )
    fig = go.Figure(data = [trace], layout = layout)
    fig.show()
    #plt.plot(fig)
```

```
In [49]: modelSummary(xTrain, yTrain, xTest, yTest)
```

[[369 0]				
[70 2]]				
	precision	recall	f1-score	support
No	0.84	1.00	0.91	369
Yes	1.00	0.03	0.05	72
micro avg	0.84	0.84	0.84	441
macro avg	0.92	0.51	0.48	441
weighted avg	0.87	0.84	0.77	441

Feature importances

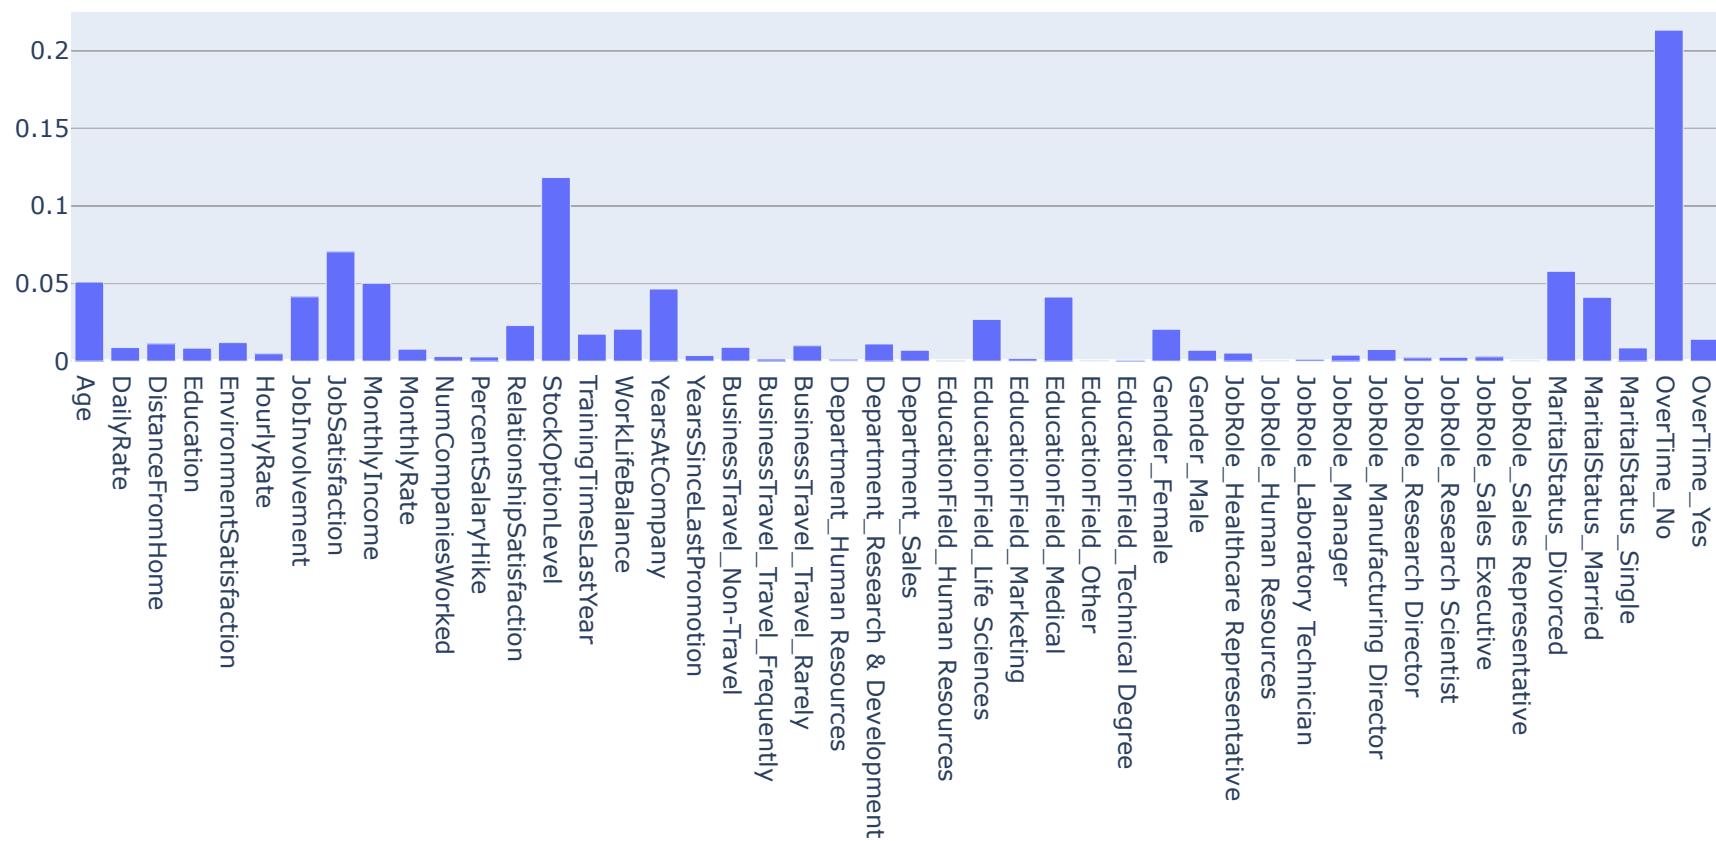


While the No-attrition class is classified correctly, the Yes-attrition class is horrible. The imbalanced set makes it difficult to build an accurate model, so we'll try resampling the training set and train again to see how the model performs and how that changes the feature importances

```
In [50]: #Fitting the model - Oversampling minority class
oversampler = SMOTE()
xOver, yOver = oversampler.fit_sample(xTrain, yTrain)
modelSummary(xOver, yOver, xTest, yTest)
```

[[338 31]	
[39 33]]	
precision recall f1-score support	
No	0.90
Yes	0.52
micro avg	0.84
macro avg	0.71
weighted avg	0.83

Feature importances

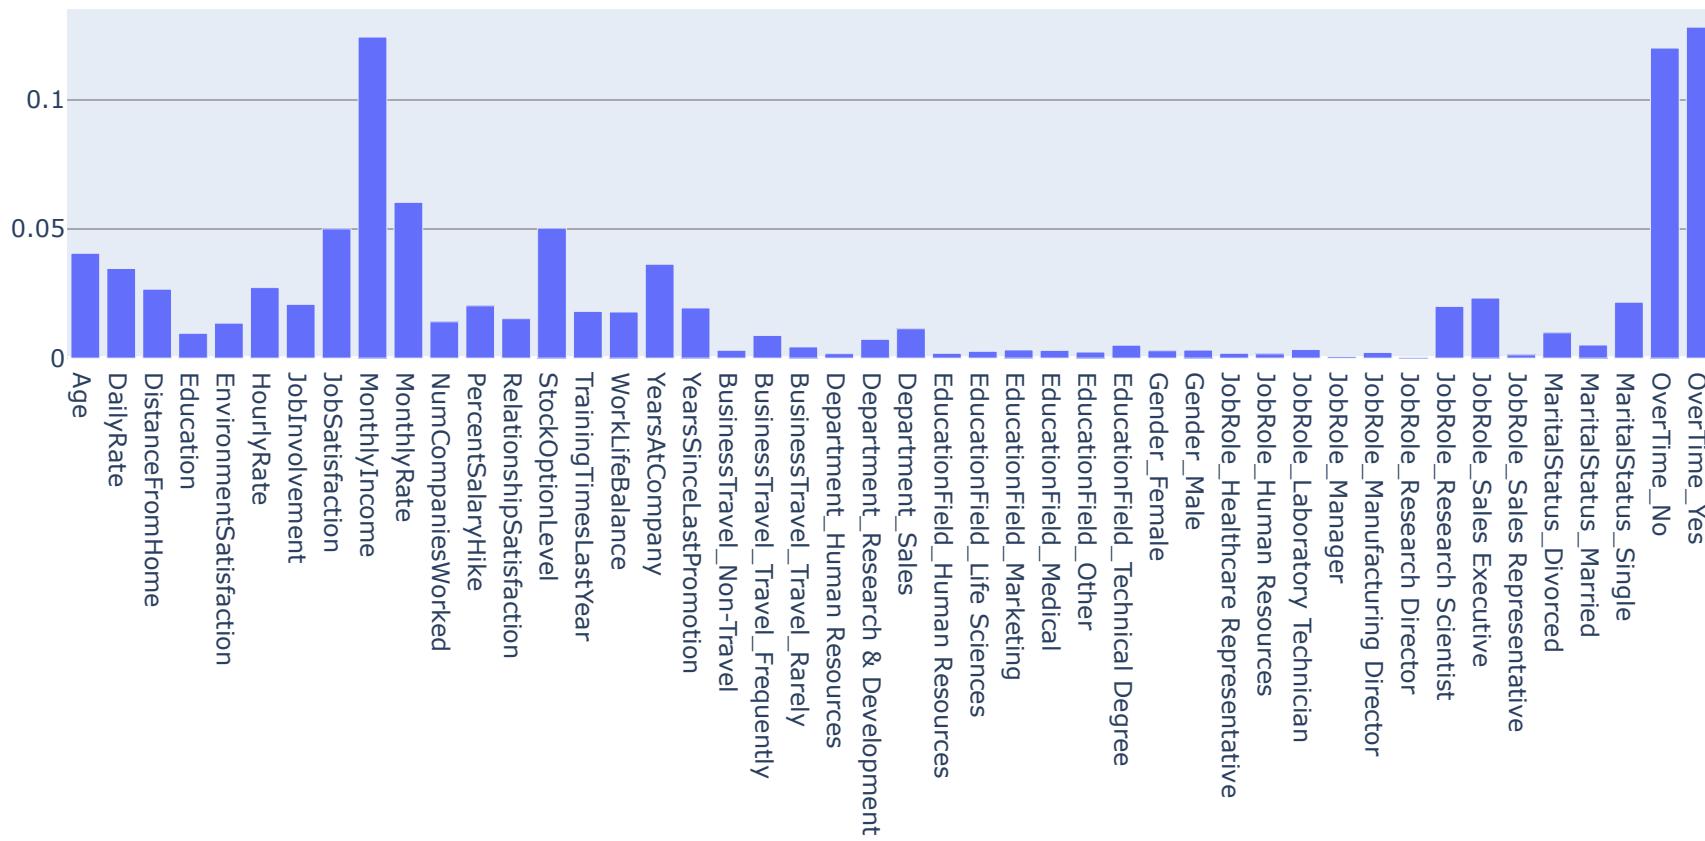


The No-attrition accuracy has dropped slightly, but in return we have a much better Yes-attrition accuracy. It's still not fantastic, but it's a lot better than before. Since this was done by oversampling the minority class, we'll try undersampling the majority class next

```
In [51]: #Fitting the model - Undersampling the majority class
undersampler = NearMiss()
xUnder, yUnder = undersampler.fit_sample(xTrain, yTrain)
modelSummary(xUnder, yUnder, xTest, yTest)
```

```
[[162 207]
 [ 18  54]]
      precision    recall   f1-score   support
No        0.90     0.44     0.59      369
Yes       0.21     0.75     0.32      72
micro avg    0.49     0.49     0.49      441
macro avg    0.55     0.59     0.46      441
weighted avg  0.79     0.49     0.55      441
```

Feature importances



Now the No-attrition class accuracy has dropped dramatically low, while the Yes-attrition is pretty decent. Overall though, it does not perform as well as the oversampled model from above

```
In [ ]:
```